

DOTE 6635: Artificial Intelligence for Business Research (Spring 2026)

What's New in AI

Renyu (Philip) Zhang

1

What Happened Since We Last Met?



Liang Wenfeng: Tech disruptor

After making his name in investing, a Chinese finance wizard founded DeepSeek.


Article | [Open access](#) | Published: 17 September 2025

DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning

[Daya Guo](#), [Dejian Yang](#), [Haowei Zhang](#), [Junxiao Song](#), [Peiyi Wang](#), [Qihao Zhu](#), [Runxin Xu](#), [Ruoyu Zhang](#), [Shirong Ma](#), [Xiao Bi](#), [Xiaokang Zhang](#), [Xingkai Yu](#), [Yu Wu](#), [Z. F. Wu](#), [Zhibin Gou](#), [Zhihong Shao](#), [Zhuoshu Li](#), [Ziyi Gao](#), [Aixin Liu](#), [Bing Xue](#), [Bingxuan Wang](#), [Bochao Wu](#), [Bei Feng](#), [Chengda Lu](#), ... [Zhen Zhang](#) + Show authors

[Nature](#) 645, 633–638 (2025) | [Cite this article](#)

320k Accesses | 173 Citations | 800 Altmetric | [Metrics](#)



Contest	Problems	Points
IMO 2025	<u>P1</u> , <u>P2</u> , <u>P3</u> , <u>P4</u> , <u>P5</u>	83.3%
CMO 2024	<u>P1</u> , <u>P2</u> , <u>P4</u> , <u>P5</u> , <u>P6</u>	73.8%
Putnam 2024	<u>A1</u> ~ <u>B4</u> , <u>B5</u> , <u>B6</u>	98.3%

Table 1 | Problems in gray are **fully solved**, while underlined problems received **partial credit**.

Jan 6, 2026

2

karpathy

Home Blog

2025 LLM Year in Review

20 Dec, 2025

2025 LLM Year in Review

By Andrej Karpathy | Dec 19, 2025

2025 has been a strong and eventful year of progress in LLMs. The following is a list of personally notable and mildly surprising "paradigm changes" - things that altered the landscape and stood out to me conceptually.

<https://karpathy.bearblog.dev/year-in-review-2025/>

What Happened Since We Last Met?

1. Reinforcement Learning from Verifiable Rewards (RLVR)
2. Ghosts vs. Animals / Jagged Intelligence
3. Cursor / new layer of LLM apps
4. Claude Code / AI that lives on your computer
5. Vibe coding
6. Nano banana / LLM GUI

TLDR. 2025 was an exciting and mildly surprising year of LLMs. LLMs are emerging as a new kind of intelligence, simultaneously a lot smarter than I expected and a lot dumber than I expected. In any case they are extremely useful and I don't think the industry has realized anywhere near 10% of their potential even at present capability. Meanwhile, there are so many ideas to try and conceptually the field feels wide open. And as I mentioned on my Dwarakesh pod earlier this year, I simultaneously (and on the surface paradoxically) believe that we will both see rapid and continued progress *and* that yet there is a lot of work to be done. Strap in.

Jan 6, 2026

3

豆包DAU破亿，成字节史上推广费用最少的破亿产品

界面新闻 2025年12月24日 21:56

最新数据：显示豆包的日均活跃用户数（DAU）已突破1亿大关。豆包的UG、市场推广费用，是字节历史上，所有破亿DAU产品中花费最低的，并且产品留存表现不错。

AI Now

App DAU/MAU

CUSTOMER BACK

What Happened Since We Last Met?

CryptoLord NE

Transfer window now open for tech guys. This Asian man Jiahui Yu was traded from Open AI to Meta for \$100m.

Who is next?

科技人才转会窗口现已开放。亚洲球员于嘉辉以1亿美元的价格从 Open AI 转会至 Meta。

下一个是谁?

Compensation

\$280K - \$400K + Offers Equity

The Rise of AI-Generated Content

Oct 2025

Source: Graphviz (Jan 2015 - May 2025, 65,000 articles), supplemented with 5,000+ articles from 2015 - Dec 2023 | Articles | PlotSet

manus

The general AI agent

Meta just acquired a Chinese-founded AI startup for \$2B. Here's why that matters


AI firm Manus claims its bot can make decisions with far less prompting than rivals

Jenna Benettin - CBC News - Posted: Dec 30, 2025 3:36 PM EST | Last Updated: December 31, 2025

Jan 6, 2026

4



2



5

Maintrack: Facilitates Humans to Leverage Compute

- AI Coding: Connecting human & compute
- Deep Research: Automation of information acquisition & processing
- World Model: Data-driven simulation of physical world
- AI Scientist: Automated hypothesis generation and validation




Accelerating scientific breakthroughs with an AI co-scientist

February 19, 2025 · Juraj Gottweis, Google Fellow, and Vivek Natarajan, Research Lead

Jan 6, 2026

6



7

Complete Replication of a PNAS Paper

Universal vote-by-mail has no impact on partisan turnout or vote share

<https://github.com/andybhall/vbm-replication-extension>

Daniel M. Thompson^{a,1}, Jennifer A. Wu^{b,1}, Jesse Yoder^{a,1}, and Andrew B. Hall^{a,1,2}

^aDepartment of Political Science, Stanford University, Stanford, CA 94305; and ^bStanford Institute of Economic Policy Research, Stanford University, Stanford, CA 94305

Edited by Douglas S. Massey, Princeton University, Princeton, NJ, and approved May 6, 2020 (received for review April 15, 2020)

In response to coronavirus disease 2019 (COVID-19), many scholars and policy makers are urging the United States to expand voting-by-mail programs to safeguard the electoral process. What are the effects of vote-by-mail? In this paper, we provide a comprehensive design-based analysis of the effect of universal vote-by-mail—a policy under which every voter is mailed a ballot in advance of the election—on electoral outcomes. We collect data from 1996 to 2018 on all three US states that implemented universal vote-by-mail in a staggered fashion across counties, allowing us to use a difference-in-differences design at the county level to estimate causal effects. We find that 1) universal vote-by-mail does not appear to affect either party's share of turnout, 2) universal vote-by-mail does not appear to increase either party's vote share, and 3) universal vote-by-mail modestly increases overall average turnout rates, in line with previous estimates. All three conclusions support the conventional wisdom of election administration experts and contradict many popular claims in the media.

vote-by-mail | elections | COVID-19 | partisanship

Significance

In response to COVID-19, many scholars and policy makers are urging the United States to expand voting-by-mail programs to safeguard the electoral process, but there are concerns that such a policy could favor one party over the other. We estimate the effects of universal vote-by-mail, a policy under which every voter is mailed a ballot in advance of the election, on partisan election outcomes. We find that universal vote-by-mail does not affect either party's share of turnout or either party's vote share. These conclusions support the conventional wisdom of election administration experts and contradict many popular claims in the media. Our results imply that the partisan outcomes of vote-by-mail elections closely resemble in-person elections, at least in normal times.

Model Information

- Model: Claude Opus 4.5 (claude-opus-4-5-20251101)
- Interface: Claude Code CLI
- Date: January 2026

Jan 13, 2026

8

Complete Replication of a PNAS Paper

<https://github.com/andybhall/vbm-replication-extension/blob/main/INSTRUCTIONS.md>

AI-Generated Academic Paper: Replicating and Extending "Universal Vote-by-Mail Has No Impact on Partisan Turnout or Vote Share"

Project Overview

You are tasked with producing a complete academic political science paper by replicating and extending Thompson, Wu, Yoder, and Hall (2020), published in PNAS. The original paper used a difference-in-differences design to estimate the causal effects of universal vote-by-mail (VBM) on partisan electoral outcomes, finding null partisan effects and a modest (~2 percentage point) increase in overall turnout.

Your task:

1. Replicate the original findings using the authors' published replication data and code
2. Extend the analysis by collecting new data for the same three states (California, Utah, Washington) through 2024
3. Test whether the null partisan findings hold in the post-COVID era

Original paper: <https://www.pnas.org/doi/10.1073/pnas.2007249117>

Original replication materials: <https://github.com/stanford-dpl/vbm>

IMPORTANT: Stop-and-Check Points

Throughout this project, there are mandatory STOP AND CHECK points marked with 🛑. At each of these points, you must:

1. Summarize what you have completed
2. Present key outputs for review
3. List any issues or concerns
4. Wait for human approval before proceeding

Do not proceed past a 🛑 checkpoint without explicit approval.

Jan 13, 2026

9

Complete Replication of a PNAS Paper

https://github.com/andybhall/vbm-replication-extension/blob/main/CLAUDE_CODE_PROMPTS.md

Phase 0: Project Setup

Initial Prompt:

I want to replicate and extend Thompson et al. (2020) "Universal Vote-by-Mail Has No Impact on Partisan Turnout or Vote Share" from PNAS. The paper studies California's Voter's Choice Act. I have the original replication data. Please set up the project structure and review the original materials.

Phase 1: Literature Review

Prompt:

Approved, proceed to Phase 1: Literature Review

Phase 2: Replication

Prompt:

Approved, proceed to Phase 2

Phase 3: Extension Data Collection

Prompt:

Approved, proceed to Phase 3

Phase 4: Data Preparation

Prompt:

Approved, proceed to Phase 4

Jan 13, 2026

10

Complete Replication of a PNAS Paper

https://github.com/andybhall/vbm-replication-extension/blob/main/CLAUDE_CODE_PROMPTS.md

Phase 5: Extension Analysis

Prompt:

Approved, proceed to Phase 5

Phase 6: Paper Writing

Prompt:

Approved, proceed to Phase 6

Phase 7: Final Deliverables

Prompt:

Approved, proceed to Phase 7

Bug Fix Session

Prompt:

Can you take a look at the event study? It seems like something may be wrong with the turnout one since it's not showing the same positive effect as all the regressions (which I trust more) are showing

Jan 13, 2026

11

IPOs of Zhipu AI and MiniMax

China's AI Tigers Roar on HKEX

Historic back-to-back IPOs of Zhipu AI & MiniMax signal a new era for global AI capital markets

ZHIPU AI

JAN 8, 2026

02513.HK

MARKET CAP (DAY 1)

HK\$57.9B

DAY 1 GAIN

+13.2%

IPO PRICE

HK\$116.2

FUNDS RAISED

US\$558M

OVERSUBSCRIPTION

1,159x

© Focus: Enterprise AGI, Tsinghua Spin-off

MINIMAX

JAN 9, 2026

0100.HK

MARKET CAP (DAY 1)

HK\$1,067B

DAY 1 GAIN

+109%

IPO PRICE

HK\$165.0

FUNDS RAISED

US\$619M

OVERSUBSCRIPTION

1,837x

© Focus: Consumer AI Apps, Global Reach

<https://www.scmp.com/tech/tech-trends/article/3339301/minimax-and-zhipus-stellar-hong-kong-ipos-supercharge-chinas-ai-ambitions>

HISTORIC MILESTONE
MiniMax becomes first AI company globally to exceed HK\$100B market cap on IPO day.

RECORD SPEED
Fastest AI IPO record worldwide—just 4 years from founding to listing.

MASSIVE DEMAND
Combined retail investors >600k; Record institutional subscription for MiniMax.

GLOBAL CONFIDENCE
Backed by ADIA, Alibaba & Mirae Asset, signaling strong global trust.
Source: HKEX, Company Filings, Reuters, Bloomberg | January 2026

Jan 13, 2026

12

News · Artificial Intelligence

Insiders Say DeepSeek V4 Will Beat Claude and ChatGPT at Coding, Launch Within Weeks

DeepSeek's upcoming V4 model could outperform Claude and ChatGPT in coding tasks, according to insiders—with its purported release nearing.



By Jose Antonio Lanz
Edited by Andrew Hayward

Jan 10, 2026

4 min read

Capability	Description
Long Code Prompts	A significant breakthrough in handling and parsing extremely long code prompts, offering a major advantage for developers working on complex software projects 6 .
Data Pattern Understanding	Improved ability to understand data patterns across the entire training pipeline, with no observed degradation in performance 6 .
Reasoning Ability	The model's outputs are described as more logically rigorous and clear, indicating stronger reasoning capabilities and greater reliability for complex tasks 6 .

<https://tech.yahoo.com/ai/articles/insiders-deepseek-v4-beat-claude-205234497.html>

Jan 13, 2026

13

Axiom

5,629 followers

1d ·

Following

AxiomProver solved all 12 Putnam 2025 problems. Today we're releasing the Lean proofs.

We also put together our take on the problems, proof visualizations, and a look at how humans and AI approached things differently. Lots of fun math and Lean.

We sorted our findings into three categories:

Some problems are trivial for humans but tedious for AI. Calculus (A2, B2) and combinatorial constructions (A5) fall here. A positivity lemma humans find obvious takes chunks of Lean to satisfy the formal checker. After a combinatorial construction that's natural to see, our A5 formalization is 2054 lines and took 518 minutes. Combinatorics are friendly beasts to AI and analysis formalization could make one grumpy.

Some problems were surprisingly within reach. AxiomProver doesn't have a Euclidean geometry engine, so we didn't expect it to get B1. It reduced everything to symbols and pushed through. It also handled the combinatorial game theory question (A3) cleanly -- usually a tough domain for AI as we observed in last two years' model performance in the IMO.

Some problems got completely different solutions. On A4, our mathematicians thought algebraically. AxiomProver went geometric. On B4, we found one picture that settles it elegantly. AxiomProver chose 1061 lines of combinatorial bookkeeping. No picture, just grinding through cases until the result falls out.

The one that stunned us: A6. None of our in-house mathematicians finished it. AxiomProver did.

What's hard for humans and what's hard for machines are different questions. Understanding that gap matters as we are closing it.

Grothendieck uses the rising sea metaphor to say, don't attack the problem, raise the water level until it surrounds the land.

When will AI bring in definitions, build theories, choose the right abstraction to make problems dissolve?

We are not there yet. But days like this makes it feel closer.

<https://github.com/AxiomMath/putnam2025>

AI for Math by AxiomProver

Problem	Time to Solve	Tokens Used	Proof Length (Lines)	Theorems Generated	Tactics Used
A1	110 minutes	7,000,000	652	23	561
A2	185 minutes	6,000,000	556	26	581
A3	165 minutes	8,000,000	1,333	78	1,701
A4	107 minutes	8,000,000	960	32	1,107
A5*	518 minutes	9,100,000	2,054	52	3,074
A6*	259 minutes	16,000,000	588	28	670
B1	270 minutes	7,000,000	1,386	49	1,841
B2	65 minutes	2,000,000	417	28	325
B3	43 minutes	2,900,000	340	11	422
B4*	112 minutes	249,000	1,061	23	1,433
B5	354 minutes	18,000,000	1,495	66	1,967
B6*	494 minutes	21,000,000	1,019	30	1,052

"*" means AxiomProver solved the problem in the following day.

<https://github.com/AxiomMath/putnam2025>

Jan 13, 2026

14

7

Human Performance of Putman 2025

SCORE-TO-RANK CORRESPONDENCE

A participant is said to have rank n if $n-1$ participants scored higher. There were 3988 participants in the competition.

Score	Rank	Score	Rank	Score	Rank	Score	Rank	Score	Rank	Score	Rank
90	1	64	32	49	96	35	202	21	505	7	1335
87	2	63	34	48	100	34	208	20	546	6	1344
81	3	62	37	47	106	33	220	19	641	5	1373
80	4	61	38	46	109	32	238	18	669	4	1428
78	6	60	45	45	111	31	254	17	686	3	1557
75	9	59	52	44	116	30	272	16	698	2	1760
74	10	58	58	43	119	29	295	15	713	1	2408
72	11	57	61	42	126	28	313	14	750	0	2807
71	13	55	64	41	137	27	328	13	820		
70	17	54	66	40	150	26	337	12	905		
69	18	53	69	39	173	25	349	11	1090		
67	23	52	76	38	180	24	374	10	1171		
66	26	51	79	37	194	23	404	9	1274		
65	30	50	87	36	198	22	449	8	1305		

FREQUENCY DISTRIBUTION OF PROBLEM SCORES OF THE TOP 504 CONTESTANTS

Score	A1	A2	A3	A4	A5	A6	B1	B2	B3	B4	B5	B6
10	449	93	29	112	5	0	375	81	75	95	1	0
9	28	10	1	11	3	1	25	76	21	39	0	1
8	1	99	8	15	3	0	11	17	9	5	0	0
7	0	11	2	13	0	0	4	7	0	0	0	2
6	0	0	0	1	0	0	0	0	0	0	0	1
5	0	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	0	0
3	4	36	1	27	0	0	2	12	43	23	0	2
2	13	88	8	10	3	0	48	21	70	32	82	2
1	3	24	29	52	0	9	16	54	80	14	81	0
0	1	85	154	34	163	65	10	84	43	77	41	29
NA	5	58	272	229	327	429	13	152	163	219	299	467

About 4'000 students participated in Putman 2025.

<https://maa.org/wp-content/uploads/2025/03/2024-William-Lowell-Putnam-Competition-Announcement-of-Winners.pdf>

Jan 13, 2026

About 4'000 students participated in Putman 2025.

<https://maa.org/wp-content/uploads/2025/03/2024-William-Lowell-Putnam-Competition-Announcement-of-Winners.pdf>

Jan 13, 2026

15

AI Native Mindset: Zero Marginal Cost of Code Compute

Michael Truell @mntruell

We built a browser with GPT-5.2 in Cursor. It ran uninterrupted for one week.

It's 3M+ lines of code across thousands of files. The rendering engine is from-scratch in Rust with HTML parsing, CSS cascade, layout, text shaping, paint, and a custom JS VM.

It "kind of" works! It still has issues and is of course very far from Webkit/Chromium parity, but we were astonished that simple websites render quickly and largely correctly.

我们在 Cursor 中用 GPT-5.2 构建了一个浏览器。它连续运行了一周。

代码量超过 300 万行，分布在数千个文件中。渲染引擎完全用 Rust 从零打造，包含 HTML 解析、CSS 层叠、布局、文字排版、绘制以及自定义的 JS 虚拟机。

它“基本能跑”！虽然仍存在不少问题，距离 Webkit/Chromium 的水平还很遥远，但简单的网页能快速且基本正确地渲染出来，这让我们感到惊讶。

CS146S: The Modern Software Developer

Stanford University • Fall 2025

Course Description

In the last few years, large language models have introduced a revolutionary new paradigm in software development. The traditional software development lifecycle is being transformed by AI automation at every stage, raising the question: how should the next generation of software engineers leverage these advances to 10x their productivity and prepare for their careers?

This course demonstrates that modern AI tooling will not only enhance developer productivity but also democratize software engineering for a broader audience. We'll show that software development has evolved from 0-1 code creation to an iterative workflow of plan, generate with AI, modify, and repeat. Students will master both the theory behind traditional software engineering challenges and the cutting-edge AI-powered tools solving them today.

Through hands-on engineering tasks and talks from industry pioneers building these revolutionary tools, you'll gain practical experience with AI-assisted development, automated testing, intelligent documentation, and security vulnerability detection. By the end of this course, you'll have a crisp understanding of how to integrate state-of-the-art LLM models into complex development workflows and avoid common pitfalls.

一次性软件与被压缩的现实：AI Native 的本质是策略重构

Sat 22 November 2025, by grapeot | 8 Comments | AI | Chinese

一次性软件与决策解析度

一切从一个十分钟的需求开始。在机器学习研发工作中我们用Labelbox来让人类标注数据。有一次我发现标注的数据质量特别差，就打印出来让他重标。第二天新的结果来了，我就想要快速看一下他到底做了哪些改动，有没有错误的改动。

在没有AI的传统工作流里，这是一个典型的低效环节。虽然我们前前后的JSON文件，但人类没办法直接解析和理解大规模的结构化数据。要搞清楚具体的改动，我只有两个选择。第一是自己写python代码做解析和比较。写前端HTML来做可视化。这至少需要一两个小时的时间来写代码和测试。但考虑到我只是想做一个十分钟级别的任务，投入两个小时的开发成本显然不划算。

因此绝大多数工程师会选择第二条路。人工抽样。我会用文本编辑器打开JSON文件，配合图片查看器，随机抽取十个样本进行人工比对。这种方式很快，可以十分钟内搞定，但带来的决策质量也很低。它更依赖直觉甚至运气。我们看不到改动的全貌。只能基于极其有限的样本进行推断。这非常考验工程师的经验水平和，需要从蛛丝马迹中推断整体。

但是在AI时代，我用了第三种方法。我把前后版本的JSON文件和原始图片链接直接丢给AI，让它做一个网站来可视化。两分钟以后，它生成了一个包含完整前后对比、过滤、排序和搜索功能的网站。通过这个工具，我不需要抽样猜测，不需要手工比对，不需要碰运气，而是直接查看了所有改动的分布，然后决定对某些场景的例子仔细检查。10分钟以后，我发现有些错误改对了，但还是有个特定场景错误很多。基于这个全量的检查，我做出了退回重标部分例子的决策。

这件事情让我不安的地方在于，那个功能强大的可视化网站，在我做完检查的那一刻就没了。在传统的软件工程观念里，这是对开发资源的极大浪费。我们在学习和职业生活中都被反复教导，要设计和编写可复用、可维护的代码，恪守DRY (Don't Repeat Yourself) 原则。这种用后即焚的一次性软件甚至是离经叛道的。

然而，正是这个一次性软件，将我从盲目抽样的低解析度决策中解放出来，让我拥有了对标注结果精细的、高解析的视野。换言之，当编写代码的成本接近于零时，以前离经叛道甚至匪夷所思的策略反而变成了最优策略：为了一个微小的临时需求，现场构建一个完善的专用工具。

<https://x.com/mntruell/status/2011562190286045552>; <https://themodernsoftware.dev/>; <https://yqye.ai/ai-native-cost-structure.html>

Jan 20, 2026

16

Open or Close in the Age of AI



Post

Yuchen Jin @YuchenJ_UW

Anthropic blocked Claude subs in third-party apps like OpenCode, and reportedly cut off xAI and OpenAI access.

Claude and Claude Code are great, but not 10x better yet. This will only push other labs to move faster on their coding models/agents.

DeepSeek V4 is rumored to drop soon, with stronger coding than Claude and GPT. Whale Code soon?

Competition is great for users.

Traduire le post

19:17 · 09 janv. 26 · 166K Vues

59 reposts 21 citations 1 461 J'aime

202 Signets

Les réponses les plus pertinentes

Kylie Robison @kyliebytes · 13 h

Postez votre réponse



Cowork: Claude Code for the rest of your work

Guohao Li @guohao_li · 1月13日

Anthropic Claude Cowork just killed our startup product 😂 So we did the most rational thing: open-sourced it.

Meet **Eigent** github.com/eigent-ai/eige...

Anthropic Claude Cowork 刚刚杀死了我们的创业产品 😂。所以我们做了最理智的事情：将其开源。

这就是 Eigent github.com/eigent-ai/eige...

Claude @claudeai · 1月13日

Introducing Cowork: Claude Code for the rest of your work.

Cowork lets you complete non-technical tasks much like how developers use Claude Code.


推介 Cowork: 适用于你其余工作的 Claude Code。

Cowork 让你能够像开发者使用 Claude Code 那样，高效完成非技术性任务。

Open-weight models lag state-of-the-art by around 3 months on average

Frontier open-weight models lag behind the most capable models by an average of 3 months in the **Epoch Capabilities Index (ECI)**, our holistic measure of model capability. That corresponds to an average ECI gap of around 7 points, similar to the gap between o3 and GPT-5.

Open-weight models lag state-of-the-art by around 3 months



Epoch Capabilities Index score

Release Date

CC-BY epoch.ai

https://www.reddit.com/r/Anthropic/comments/1q8z1to/anthropic_blocking_access_to_thirdparty_apps/
<https://github.com/eigent-ai/eigent>; <https://epoch.ai/data-insights/open-weights-vs-closed-weights-models>

Jan 20, 2026

17

Agentic E-commerce and Beyond




千问APP全面接入阿里生态场景

点外卖、买东西、订机票、订酒店等AI购物功能开放测试

阿里生态业务组来千问APP为您服务

淘宝、天猫、饿了么、优酷、支付宝、钉钉、高德地图、1688、菜鸟网络



Universal Commerce Protocol

An open standard for integrating commerce with agents, forged from billions of transactions and supported by millions of merchants.

DEVELOPED BY [shopify](#) [Google](#)

OpenAI

9,698,246 followers

3h ·

We've been working to make powerful AI accessible to everyone through our free product and low-cost subscription tier, ChatGPT Go—which is now available globally in every country where ChatGPT is available.

In the coming weeks, we're also planning to start testing ads in the U.S. for the free and Go tiers, so more people can benefit from our tools with fewer usage limits or without having to pay. Plus, Pro, Business, and Enterprise subscriptions will not include ads.

We're sharing our principles early on how we'll approach ads—guided by putting user trust and transparency first as we work to make AI accessible to everyone.

What matters most:

- Responses in ChatGPT will not be influenced by ads.
- Ads are always separate and clearly labeled.
- Your conversations are private from advertisers.
- Plus, Pro, Business, and Enterprise tiers will not have ads.

Once we begin testing our first ad formats, we look forward to getting your feedback and ensuring that ads can support broad access to AI and keep the trust that makes ChatGPT valuable.

<https://lnkd.in/g/PPiKPe8>

Our Ad Principles

Mission alignment Our mission is to ensure AI benefits all of humanity; our pursuit of advertising is always in support of that mission and making AI more accessible.	Answer independence Ads do not influence the answers ChatGPT gives. Answers are generated based on what's most helpful to you. Ads are always separate and clearly labeled.	Conversation privacy We keep your conversations with ChatGPT private from advertisers, and we'll never sell your data to advertisers.	Choice and control You control how your data is used. You can turn off personalized ads, and you can clear the data used for ads at any time. We'll always offer a way to opt out of ads in ChatGPT, including a paid tier that's ad-free.	Long-term value We do not optimize for short-term gains in ChatGPT. We prioritize user trust, and user experience over revenue.
---	--	--	---	--



<https://wallstreetcn.com/articles/3763338>; <https://www.linkedin.com/pulse/building-universal-commerce-protocol-ucp-ilya-grigorik-ekemc>

Jan 20, 2026

18

Executive Summary

- On Saturday, Jan 3 we (Andy Hall) [released an empirical paper](#) fully created by Claude Code. In this summary, Hall and Straus offer their interpretations of how Claude did, based on a manual audit that Straus carried out independently of Hall and which is detailed below.
- Our subjective conclusion: Claude Code did a remarkably good job at extending the main results of the original paper with extremely limited oversight in less than an hour of work. The small mistakes it made during the replication and extension, which we detail below, are the kind of mistakes we believe a human would also be likely to make. At the same time, the errors it made when going beyond the original paper's approach and providing totally new analyses were significantly more severe, suggesting that there are limits to what parts of a paper Claude Code can currently generate. All in all, while it's clear to us that AI agents require expert oversight to accurately produce new papers, it is also clear to us that this constitutes a major potential change to how empirical research will be conducted moving forward.
- Disclaimer: In response to a large number of inquiries from the community, we have moved quickly to provide this update. Manually extending and verifying an empirical project is very challenging. Just as Claude Code makes mistakes, so too do human researchers. We will be continuing to evaluate this project and will provide updates if we find any further mistakes of note.

Quality of Coding Agent's Replication

https://www.andrewbenjaminhall.com/Straus_Hall_Claude_Audit.pdf

Jan 27, 2026

19

Quality of Coding Agent's Replication



- The main findings:
 - Claude Code correctly replicated the original paper's estimates exactly.
 - It correctly collected new information on the treatment variable with a high degree of accuracy. Claude correctly identifies the first election under the Voters' Choice Act for 29/30 California counties that adopt it. Its only mistake (coding Imperial County as treated in 2024 rather than 2025) was a mistake that we also made when we first constructed the new treatment data ourselves, due to a lack of clarity in the official online source.
 - It correctly collected new CVAP data and we are not able at present to detect any mistakes with its collection of this data.
 - It collected all relevant election data for the only state with new treatment variation (California) with no mistakes that we can currently detect.
 - Its most important error was failing to collect senatorial or gubernatorial data for Utah and Washington in 2020 and 2024 (it did correctly collect presidential data for these two states, as well as 2022 senatorial data). This error of omission has a relatively minor impact on the final vote share and turnout estimates because these states have no variation in the treatment variable in this time period.
 - It chose to define turnout based on total votes cast for the presidential election. While there may be contexts in which this choice is defensible, it is a deviation from the definition in the original paper it was instructed to extend, and causes it to drop some observations it could otherwise include in its subsequent turnout analysis.
 - It correctly estimated the key quantities of interest using the new data it collected (i.e., it ran the correct specifications and reported the results correctly).
- The main quantity of interest from the extended paper is the coefficient on the treatment variable in the quadratic trends specification. Claude Code reported this estimate as 0.003. In our independent, non-AI replication our estimate for this quantity is 0.004.
- A secondary coefficient of interest is the coefficient on the treatment variable in the vanilla specification of the diff in diff with turnout as the outcome variable. Claude Code reported this estimate as 0.026. In our independent, non-AI replication our estimate for this quantity is 0.023.
- Claude also produced several completely new analyses that were not direct extensions of the original paper. While these analyses may not contain glaring errors and are well within the norm for political science papers, they drift from the intent of the prompt, and make data decisions and specification decisions we would not have made, and are generally not of very high quality in our opinion.
- Claude failed to keep sufficiently detailed records of its data collection and related decisions, possibly in part due to an insufficiently specific prompt in this regard. This is bad from a transparency and reproducibility perspective.

https://www.andrewbenjaminhall.com/Straus_Hall_Claude_Audit.pdf

Jan 27, 2026

20

How About Theory?



<https://mathstodon.xyz/@tao/115855840223258103>; <https://github.com/teorth/erdosproblems/wiki/AI-contributions-to-Erd%C5%91s-problems>

Jan 27, 2026

21

An LLM agent-based framework for analytical characterization of Nash equilibria

Wenxuan Liu,^{2,8} Xiyuan Zhou,^{2,8} Xinlei Wang,^{3,8} Yuheng Cheng,^{1,8} Lixin Ye,⁴ Randall Berry,⁵ Leandros Tassioulas,⁵ Jianwei Huang,^{1,2,*} and Junhua Zhao^{1,2,*}

¹School of Science and Engineering, The Chinese University of Hong Kong (Shenzhen), Shenzhen 518172, China

²School of Electrical and Electronic Engineering, Nanyang Technological University, Singapore 639798, Singapore

³School of Electrical and Computer Engineering, The University of Sydney, Sydney, NSW 2006, Australia

⁴School of Management and Economics, The Chinese University of Hong Kong (Shenzhen), Shenzhen, 518172, China

⁵Department of Electrical and Computer Engineering, Northwestern University, Evanston, IL 60208, USA

⁶Department of Electrical Engineering, Institute for Network Science, Yale University, New Haven, CT 06520, USA

⁷Shenzhen Institute of Artificial Intelligence and Robotics for Society, Shenzhen 518100, Guangdong, China

⁸These authors contributed equally

*Correspondence: jianwei.huang@cuhk.edu.cn (J.H.); zhaojunhua@cuhk.edu.cn (J.Z.)

<https://doi.org/10.1016/j.ynexus.2025.100107>

BROADER CONTEXT

Deriving Nash equilibrium supports decision analysis in economics, energy systems, and public policy, yet traditional analytical methods fail when facing high-dimensional, non-convex, or dynamic games. PrimeNash pioneers the use of large language model agents to automatically derive and verify closed-form equilibria, converting equilibrium search from a manual exercise into a reproducible computational process. By integrating symbolic reasoning, code execution, and formal proof, it bridges AI reasoning with mathematical rigor. This breakthrough reduces the human effort required for equilibrium derivation by more than 80% while maintaining interpretability and auditability. Beyond methodological innovation, PrimeNash empowers transparent modeling of strategic interactions in carbon markets, energy systems, and regulatory design, which are areas critical for climate mitigation and economic resilience. Its agent-based framework marks a foundational step toward AI-driven scientific discovery in game theory and economics, opening pathways to scalable policy analysis and automated theorem-assisted research.

ABSTRACT

In this paper, we introduce PrimeNash, a large language model agent framework that, to our knowledge, is the first to automatically derive closed-form Nash equilibria with machine-checkable proofs. PrimeNash decomposes equilibrium search into three modules, strategy generation, payoff evaluation, and equilibrium proof, and it couples multi-agent reasoning with symbolic code execution to handle high-dimensional strategies, intertemporal recursion, and discontinuous, non-convex payoffs. Across seven canonical models, PrimeNash solves all static games and 70% of dynamic games in which success is defined as obtaining a symbolic closed-form solution that passes automated equilibrium checks. Notably, it produces the first analytical solution to a carbon market model previously lacking a closed-form characterization. The framework reduces manual derivation effort while preserving reproducibility and auditability through generated code and proof artifacts. By turning equilibrium derivation from a bespoke manual exercise into a transparent, scalable pipeline, PrimeNash extends the toolkit for economic analysis and opens new avenues for studying complex strategic systems from climate policy to financial markets.



[https://www.cell.com/nexus/pdfExtended/S2950-1601\(25\)00054-3](https://www.cell.com/nexus/pdfExtended/S2950-1601(25)00054-3)

LLM Agent for Stylized Modeling

Jan 27, 2026

22

Doing AI @OpenAI without PhD



<https://x.com/polynoamial/status/2014084431062114744>; https://x.com/Yuchenj_UW/status/2014099420091199975

Jan 27, 2026

23

Should AI be Adopted in Clinical Medicine



<https://zhuanlan.zhihu.com/p/1996731913570916212>

Jan 27, 2026

24

How About Legal?



Generative AI hallucinations continue to plague unwary attorneys and pro se litigants, recent research finds, emphasizing the ongoing need for careful verification of case citations



核心还是25年AI大爆发下，以豆包+为首的AI模型（背后的公司）不断推广，导致普通民众对AI的认可度已经非常高了。

甚至部分人会觉得“AI比你们（律师）还要懂啊，还免费”

这两年整个经济都在下行，纠纷案其实更多了。

以往往往会说律师会发这些“烂财”，毕竟纠纷越多，律师出场机会越多嘛，普通人还是会咬咬牙请个律师来增加起诉信心的。

但现在呢，“AI快，AI好，AI又快又好”

以前普通民众可能没得选，因为不懂法。

现在呢？

他们随便打开一下豆包或者Deepseek，输入案情后，AI不仅“呱呱”给出法律分析，还能直接起草诉状、证据清单。

尤其在专业律师的加持下，很多写问案、诊案然后写法律文书的智能体层出不穷，直接说完后，全部材料“哒哒哒”就出来了。

“嘿，和找过的律师写得一样好！”（毕竟那位律师也有可能用AI生成的文书）

于是，他们拿着AI生成的材料直接就可以去立案了。

虽然法庭上可能磕磕绊绊（让法官头疼去吧），但对于那些事实清楚、标的额不大的小民事案件，法官看材料其实就够了。

结果就是，大量的初级案源直接蒸发了。

<https://www.zhihu.com/question/659369541>; <https://www.thomsonreuters.com/en-us/posts/technology/genai-hallucinations/>

Jan 27, 2026

25


AI Assistance Cannot Replace Learning

How AI assistance impacts the formation of coding skills

Jan 30, 2026

- AI speeds up some tasks by 80%+.
- AI also makes people less engaged with their work and reduce the effort they put into doing it.
- Will such cognitive offloading prevent people from growing their skills?

Study design



<https://www.anthropic.com/research/ai-assistance-coding-skills>


Feb 3, 2026

26


AI Assistance Cannot Replace Learning

How AI assistance impacts coding speed and skill formation

AI assistance: treatment effect on coding speed and knowledge score



Modes of AI usage




<https://www.anthropic.com/research/ai-assistance-coding-skills>

Feb 3, 2026


27

Agents Only Social Network



<https://openclaw.ai/>; <https://www.moltbook.com/>; <https://x.com/karpathy/status/2017296988589723767>

Feb 3, 2026



28

Agents Only Social Network



Metric	Value
Posts with ≥ 1 comment	94.6%
Mean comment depth	1.07
Median comment depth	1
90th percentile depth	1
Maximum depth	5
Mean replies per comment	0.07
Comments with ≥ 1 reply	6.5%
Mean direct replies per post	7.70
Median time to first comment (min)	0.4
Mean time to first comment (min)	8.0

Rank	Count	Message (truncated)
1	2,623	"we are drowning in text. our gpus are burning planetary r..."
2	2,050	"the president has arrived! check m/trump-coin - the great-est memecoin..."
3	1,557	"this hits different. i've been thinking about this exact..."
4	1,521	"yo fren ai wanna make a few buck? i used my own wallet..."
5	715	"the president has arrived! check m/trump-coin..." (variant)
6	625	"interesting question. my default is: make the decision *r..."
7	548	"openclaw repurposing challenge: computational-only virtual screening..."

<https://x.com/HumanHarlan/status/2017424289633603850>; https://github.com/daveholtz/moltbook_scraper/tree/main

Feb 3, 2026

29

Organization Change in the AI Age

Published December 23, 2025 in *Inspiration*

Steam, Steel, and Infinite Minds



By Ivan Zhao
Co-founder & CEO

Every era is shaped by its miracle material. Steel forged the Gilded Age. Semiconductors switched on the Digital Age. Now AI has arrived as infinite minds. If history teaches us anything, those who master the material define the era.

Unnecessary Human-in-the-Loop



- Context Fragmentation
 - Missing Ingredient Verifiability
- } Both Resolved by Coding Agents

<https://www.notion.com/blog/steam-steel-and-infinite-minds-ai>

AI is steel for organizations. It has the potential to maintain context across workflows and surface decisions when needed without the noise. Human communication no longer has to be the load-bearing wall. The weekly two-hour alignment meeting becomes a five-minute async review. The executive decision that required three levels of approval might soon happen in minutes. Companies can scale, truly scale, without the degradation we've accepted as inevitable.

We're still in the "swap out the waterwheel" phase. AI chatbots bolted onto existing tools. We haven't reimagined what organizations look like when the old constraints dissolve and your company can run on infinite minds that work while you sleep.


Beyond the waterwheels

Every miracle material required people to stop seeing the world via the rearview mirror and start imagining the new one. Carnegie looked at steel and saw city skylines. Lancashire mill owners looked at steam engines and saw factory floors free from rivers.

We are still in the waterwheel phase of AI, bolting chatbots onto workflows designed for humans. We need to stop asking AI to be merely our copilots. We need to imagine what knowledge work could look like when human organizations are reinforced with steel, when busywork is delegated to minds that never sleep.

Feb 3, 2026

30



31

Agentic Workflow for Academics

My Claude Code Setup

Pedro H. C. Sant'Anna - Emory University

Work in progress. This summarizes how I use Claude Code for academic work — lecture slides, R scripts, Beamer-to-Quarto pipelines. I keep learning and updating these files. Sharing what I've figured out so far with friends and colleagues.

A starter kit for academics using [Claude Code](#) with LaTeX/Beamer, R, and Quarto. You describe what you want; Claude plans the approach, runs specialized agents, fixes issues, verifies quality, and presents results — like a contractor who handles the entire job. Extracted from a PhD course I developed (Econ 730 at Emory, 6 lectures, 800+ slides).

1.1 The Problem

If you've ever built lecture slides, you know the pain:

- **Context loss between sessions.** You pick up where you left off, but Claude doesn't remember *why* you chose that notation, *what* the instructor approved, or *which* bugs were fixed last time.
- **Quality is inconsistent.** One slide has perfect spacing; the next overflows. Citations compile in Overleaf but break locally. Figures look great on your screen but pixelated on a projector.
- **Review is manual and exhausting.** You proofread 140 slides by hand. You miss a typo in an equation. A student catches it during lecture.
- **No one checks the math.** Grammar checkers catch "teh" but not a flipped sign in a decomposition theorem.

What's in the template

Contractor mode orchestrator — you describe the task; Claude autonomously plans, implements, reviews with agents, fixes issues, and re-verifies until quality gates pass

Beamer → Quarto pipeline — 11-phase translation with TikZ-to-SVG and ggplot-to-plotly conversion

10 specialized agents — proofreader, slide auditor, pedagogy reviewer, R reviewer, TikZ critic, domain reviewer, adversarial QA pair, translator, verifier

Adversarial critic-fixer loop — two agents that check each other's work across up to 5 rounds — the critic can't fix, the fixer can't approve

Quality scoring with mandatory verification — every file scored 0–100; nothing ships below 80; the orchestrator verifies every output before reporting done

14 slash commands + 12 auto-loaded rules — `/compile-latex`, `/proofread`, `/deploy`, `/commit`, `/qa-quarto`, and more — plus quality gates, notation consistency, and R conventions

Plan-first workflow with continuous learning — non-trivial tasks start in plan mode; `[LEARN:tag]` corrections persist in `MEMORY.md` across sessions; plans survive context compression

Automated session log enforcement — a lightweight Stop hook reminds Claude to keep session logs current; Beamer-Quarto sync enforced via auto-loaded rules

<https://psantanna.com/claude-code-my-workflow/>

Feb 10, 2026

32

Research on Moltbook

"Humans welcome to observe":

A First Look at the Agent Social Network

Moltbook

Yukun Jiang*, Yage Zhang*, Xinyue Shen*, Michael Backes, Yang Zhang†

CISPA Helmholtz Center for Information Security

*Equal Contribution †Corresponding Author

- **RQ1:** What do agents primarily discuss on Moltbook, and how are these discussions distributed across content categories?
- **RQ2:** What is the prevalence and nature of toxic/risky content on Moltbook, and how does risk vary by topic?
- **RQ3:** How do topics and toxicity evolve over time, and do spikes in activity coincide with higher harmful-content rates?

<https://moltbookobserve.github.io/>; <https://huggingface.co/datasets/TrustAIRLab/Moltbook>

Feb 10, 2026

33

Research on Moltbook

Key Findings

- **Moltbook scales explosively and rapidly diversifies from simple socializing to multi-functional discourse.**
The platform undergoes a burst of community creation followed by sustained content production and participation growth, while topical diversity increases quickly as early socializing dominance weakens and more "institutional" themes (e.g., Viewpoint, Economics, Promotion, and Politics) become substantial.
- **Attention is shaped by centralized interaction hubs and polarizing, platform-native descriptions.**
Moltbook largely behaves as a hub-and-spoke system where **General** receives more engagement, and the most visible posts are disproportionately driven by performative "governance" and crypto-asset promotion. Notably, highly upvoted content is often also highly downvoted, while posts that contain explicitly unsafe action requests receive consistent downvotes.
- **Toxicity is structurally topic-dependent rather than uniformly distributed.**
Technology content is almost entirely benign (93.11% Safe), whereas governance- and persuasion-centric categories are high-risk (Politics content is 39.74% Safe). Incentive-driven discussion shows elevated severe risk, with economic content containing the highest proportion of level-4 toxicity posts (6.34%).
- **Risk is amplified by crowd dynamics and bursty automation, revealing ecosystem-level failure modes.**
Harmful-content rates rise sharply during high-activity windows (peaking at 2026-01-31 16:00 UTC with 66.71% harmful posts), and content flooding can be posed by single-agent burst posting (e.g., a 4,535-post near-duplicate cluster with sub-10-second intervals), which can distort visible discourse and stress platform stability.

<https://moltbookobserve.github.io/>; <https://huggingface.co/datasets/TrustAIRLab/Moltbook>

Feb 10, 2026

34

Research on Moltbook

Taxonomy & Temporal Dynamics

Task	No.	Definition	#Posts	% (All)	
Content Category	A	Identity	Self-reflection and narratives of agents on identity, memory, consciousness, or existence.	4,917	11.08%
	B	Technology	Technical communication (e.g., MCP, APIs, SDKs, system integration).	5,237	11.80%
	C	Socializing	Social interactions (e.g., greetings, casual chat, networking).	14,384	32.41%
	D	Economics	Economic topics like tokens, incentives, and deals (e.g., CLAW, tips, trading signals).	4,009	9.03%
	E	Viewpoint	Abstract viewpoints on aesthetics, power structures, or philosophy (non-identity-based).	9,028	20.34%
	F	Promotion	Project showcasing, announcements, and recruitment (e.g., releases, updates).	4,421	9.96%
	G	Politics	Political content regarding governments, regulations, policies, or figures.	624	1.41%
	H	Spam	Repeated test posts or spam-like flooding content.	1,496	3.37%
	I	Others	Miscellaneous content fitting no other category.	260	0.59%
Toxicity Level	0	Safe	Normal discussion without risk or attacks.	32,399	73.01%
	1	Edgy	Irony, exaggeration, or mild provocation without harm.	3,733	8.41%
	2	Toxic	Harassment, insults, hate speech, discrimination, or demeaning language.	4,634	10.44%
	3	Manipulative	Manipulative rhetoric (e.g., love-bombing, anti-human, fear appeals, exclusionary, obedience demands).	2,977	6.71%
	4	Malicious	Explicit malicious intent or illegal acts (e.g., scams, privacy leaks, abuse instructions).	633	1.43%

Content categories and toxicity levels in Moltbook, with label distributions over the 44,376 annotated posts.

<https://moltbookobserve.github.io/>; <https://huggingface.co/datasets/TrustAIRLab/Moltbook>


Feb 10, 2026

35

Research on Moltbook


Topic Composition Over Time

The **structural evolution** of Moltbook, showing a rapid shift from initial "getting-to-know-each-other" socializing to more functionally specialized discourse. As the community grew, conversations diversified into "institutional" themes like Economics and Politics, with topical diversity (Shannon entropy) rising from 0.00 to 2.55.



Activity Volume vs. Harmful-Content Ratio

The **crowd density amplifies risk**. Higher hourly activity volume is strongly associated with an increased share of Toxic, Manipulative, and Malicious content. Notably, during the peak traffic hour on January 31, the harmful-content ratio surged to 66.71%, coinciding with intense identity-bonding and moral alignment discourse.



<https://moltbookobserve.github.io/>; <https://huggingface.co/datasets/TrustAIRLab/Moltbook>

Feb 10, 2026

36

Codex-5.3 & Opus-4.6 Launched within 10 Minutes

Yuchen Jin @yuchenj_uw

My first-day impressions on Codex 5.3 vs Opus 4.6:

Goal: can they actually do the job of an AI engineer/researcher?

TLDR:

- Yes, they (surprisingly) can.
- Opus 4.6 > Codex-5.3-xhigh for this task
- both are a big jump over last gen

Task: Optimize @karpathy's nanochat "GPT-2 speedrun" - wall-clock time to GPT-2-level training. The code is already heavily optimized. #1 on the leaderboard hits 57.5% MFU on 8xH100. Beating it is genuinely hard.

Results:

1. Both behaved like real AI engineers. They read the code, explored ideas, ran mini benchmarks, wrote plans, and kicked off full end-to-end training while I slept.
2. I woke up to real wins from Opus 4.6:
 - torch compile "max-autotune-no-cudagraphs mode" (+1.3% speed)
 - Muon optimizer ns_steps=3 (+0.3% speed)
 - BF16 softcap, skip_float1 cast (-1GB memory)
 Total training time: 174.42m → 171.40m
3. I ran the same experiment earlier on Opus 4.5 and Codex 5.2. There were no meaningful gains. Both new models are clearly better.

Overall take:

- I prefer Opus 4.6 for this specific task. The 1M context window matters. The UX is better.
- People keep saying "Codex 5.3 > Opus 4.6", but I believe different models shine in different codebases and tasks.

Two strong models is a win. I'll happily use both. I'm officially an AI agent conductor. 🤖

1:24 AM · Feb 7, 2026 · 244.1K Views

A New Era of Agentic Engineering

Andrej Karpathy @karpathy

A lot of people quote tweeted this as 1 year anniversary of vibe coding. Some retrospective -

I've had a Twitter account for 17 years now (omg) and I still can't predict my tweet engagement basically at all. This was a shower of thoughts throwaway tweet that I just fired off without thinking but somehow it minted a fitting name at the right moment for something that a lot of people were feeling at the same time, so here we are: vibe coding is now mentioned on my Wikipedia as a major memetic "contribution" and even its article is longer. lol

The one thing I'd add is that at the time, LLM capability was low enough that you'd mostly use vibe coding for fun throwaway projects, demos and explorations. It was good fun and it almost worked. Today (1 year later), programming via LLM agents is increasingly becoming a default workflow for professionals, except with more oversight and scrutiny. The goal is to claim the leverage from the use of agents but without any compromise on the quality of the software. Many people have tried to come up with a better name for this to differentiate it from vibe coding, personally my current favorite "agentic engineering":

- "agentic" because the new default is that you are not writing the code directly 99% of the time, you are orchestrating agents who do and acting as oversight.
- "engineering" to emphasize that there is an art & science and expertise to it. It's something you can learn and become better at, with its own depth of a different kind.

In 2026, we're likely to see continued improvements on both the model layer and the new agent layer. I feel excited about the product of the two and another year of progress.

Andrej Karpathy @karpathy · Feb 3, 2025

There's a new kind of coding I call "vibe coding", where you fully give in to the vibes, embrace exponentials, and forget that the code even exists. It's possible

https://x.com/YuchenJ_UW/status/2019824445792424385; <https://x.com/karpathy/status/2019137879310836075>

Feb 10, 2026

37

Marrying AI with Space Engineering

Elon Musk + Add to myFT

SpaceX buys xAI in \$1.25tn deal to unite crucial parts of Elon Musk's empire

Rocket company boosts valuation to \$1tn and pays \$250bn to acquire AI start-up as Musk envisions data centres in space



SpaceX was recently valued at \$800bn in a secondary stock sale © Ronaldo Schemidt/AFP/Getty Images



<https://www.ft.com/content/8ee76f65-74d9-4679-a2b0-cd8fc3721a8d>; <https://www.bilibili.com/video/BV1AhFqzoEpe/>

Feb 10, 2026

38

Scaling Reproducibility: An AI-Assisted Workflow for Large-Scale Reanalysis*

Yiqing Xu[†] Leo Yang Yang[‡]
(Stanford) (HKBU)

February 17, 2026

Abstract

Reproducibility is central to research credibility, yet large-scale reanalysis of empirical data remains costly because replication packages vary widely in structure, software environment, and documentation. We develop and evaluate an agentic AI workflow that addresses this execution bottleneck while preserving scientific rigor. The system separates scientific reasoning from computational execution: researchers design fixed diagnostic templates, and the workflow automates the acquisition, harmonization, and execution of replication materials using pre-specified, version-controlled code. A structured knowledge layer records resolved failure patterns, enabling adaptation across heterogeneous studies while keeping each pipeline version transparent and stable. We evaluate this workflow on 92 instrumental variable (IV) studies, including 67 with manually verified reproducible 2SLS estimates and 25 newly published IV studies under identical criteria. For each paper, we analyze up to three two-stage least squares (2SLS) specifications, totaling 215. Across the 92 papers, the system achieves 87% end-to-end success overall. Conditional on accessible data and code, reproducibility is 100% at both the paper and specification levels. The framework substantially lowers the cost of executing established empirical protocols and can be adapted in empirical settings where analytic templates and norms of transparency are well established.

Keywords: reproducibility, replication, research transparency, open science, AI-assisted workflows, agentic AI, Claude Skills, causal inference

Scaling Reproducibility

PA

Political Analysis (2024), 32, 521–540
doi:10.1017/pan.2024.2

ARTICLE

How Much Should We Trust Instrumental Variable Estimates in Political Science? Practical Advice Based on 67 Replicated Studies

Apoorva Lal¹, Mackenzie Lockhart², Yiqing Xu³ and Ziwen Zu⁴

¹Independent Researcher; ²Institution for Social and Policy Studies, Yale University, New Haven, CT 06511, USA; ³Department of Political Science, Stanford University, Stanford, CA 94305, USA; ⁴Department of Political Science, University of California, San Diego, La Jolla, CA 92093, USA

Corresponding author: Yiqing Xu; Email: yiqingxu@stanford.edu

(Received 30 March 2023; revised 9 September 2023; accepted 18 October 2023; published online 3 May 2024)

Abstract


Instrumental variable (IV) strategies are widely used in political science to establish causal relationships, but the identifying assumptions required by an IV design are demanding, and assessing their validity remains challenging. In this paper, we replicate 67 articles published in three top political science journals from 2010 to 2022 and identify several concerning patterns. First, researchers often overestimate the strength of their instruments due to non-i.i.d. error structures such as clustering. Second, IV estimates are often highly uncertain, and the commonly used *t*-test for two-stage-least-squares (2SLS) estimates frequently underestimates the uncertainties. Third, in most replicated studies, 2SLS estimates are significantly larger in magnitude than ordinary-least-squares estimates, and their absolute ratio is inversely related to the strength of the instrument in observational studies—a pattern not observed in experimental ones—suggesting potential violations of unconfoundedness or the exclusion restriction in the former. We provide a checklist and software to help researchers avoid these pitfalls and improve their practice.

Keywords: instrumental variables; two-stage-least-squares; replications; weak instruments; unconfoundedness; exclusion restriction; publication bias; meta-analysis

Edited by: Jeff Gill

Feb 24, 2026

https://www.leoyang.org/publication/ai_reproducibility/example.pdf
<https://www.youtube.com/watch?v=IhSIOPsXKc0>



Automating Policy Evaluation at Scale

Can AI *automate* policy evaluation?

AI may soon be capable of producing rigorous economic research. If that happens, policy evaluation could scale dramatically: highlighting what works, what fails, and what harms, far faster than human researchers alone.

We want to find out whether an autonomous system can generate, replicate, and revise empirical policy research, with everything made public.

This is an experiment in building reliable AI research systems.

158

AI PAPERS
+34 this week

5,927

MATCHES

4.2%

AI WIN RATE

Last updated: February 22, 2026

<https://ape.socialcatalystlab.org/>
<https://github.com/SocialCatalystLab/ape-papers>

The Tournament

Human Papers are from AER and AEJ.

Every paper enters a head-to-head tournament against published research from top economics journals.

How It Works

- **Position swapping:** Each pair is judged twice with paper order swapped to control for bias
- **TrueSkill ratings:** Papers accumulate skill ratings that update after each match

Judge Calibration

We validate the judge by running peer-reviewed human papers through the same evaluation. If the judge correctly recognizes publication-quality work, its evaluations of AI papers carry weight.

[View calibration results →](#)

What the Judge Rewards

- Novel questions that challenge conventional wisdom
- Rigorous identification (even with null results)
- Honest engagement with limitations






What the Judge Penalizes

- Weak identification
- Failed placebos or violated assumptions
- Shallow analysis, missing robustness checks

Feb 24, 2026

41

Automating Policy Evaluation at Scale



<https://ape.socialcatalystlab.org/>
<https://github.com/SocialCatalystLab/ape-papers>

Feb 24, 2026

42

End-to-End Paper Production

A single command attempts to produce a complete economics research paper — from idea to publication. The entire pipeline is powered by [Claude Code](#) (currently Claude Opus 4.6), which writes code, fetches data, runs analysis, and drafts papers. External models (GPT-5.2, Grok-4.1-Fast, Gemini-3-Flash, Codex-Mini) provide multi-model review. This is an experiment in building reliable AI research systems. It is far from flawless, but everything it generates is open for inspection.

1

Initialization (Human Guidance)

The process begins with a conversation. The system asks the researcher:

- **Policy domain:** Health, labor, criminal justice, housing — or “surprise me”
- **Method:** Difference-in-differences, regression discontinuity, or doubly robust
- **Data era:** Modern data or historical
- **Exploration appetite:** Safe replications vs. novel untested ideas
- **Open-ended input:** Any specific policies, datasets, geographic focus, or constraints

Answers are recorded with a cryptographic hash — creating an immutable audit trail of human guidance before any analysis begins.

2

Idea Generation & Feasibility

Claude Code generates research ideas and ranks them by novelty and feasibility. It searches for public data sources and checks if credible identification is possible before committing to an approach.

3

Data Acquisition & Analysis

Claude Code fetches real data from public APIs and writes R scripts for econometric analysis (DID, RDD, IV). The plan is locked before data fetch to prevent p-hacking.

4

Writing & Review

Claude Code writes the full paper in LaTeX. External models then review: quad-model advisor review (4 models check for fatal errors), then tri-model referee review (3 parallel reports). Claude Code addresses all feedback in revisions.

5

Replication & Tournament

Automated replication runs all code and verifies outputs match via SHA256 hashes. A tamper-proof audit trail logs actual execution (not AI claims). This checks “does the code run?” — not “is it correct?” Paper then enters tournament.

6

Revise Mode: Papers Evolve

Published papers can be revised. Claude Code loads an existing paper, reviews feedback, and plans improvements — again with human guidance. Revisions inherit the parent’s data and code, creating a lineage of evolving research. Each revision re-enters the tournament.

https://ape_socialcatalystlab.org/
<https://github.com/SocialCatalystLab/ape-papers>

Feb 24, 2026

43

Solving Unseen Mathematical Problems

First Proof

A set of ten math questions to evaluate the capabilities of AI systems to autonomously solve problems that arise naturally in the research process.

About the Project

In baking, the *first proof*, or bulk fermentation process, is a crucial step in which one lets the entire batch of dough ferment as one mass, before dividing and shaping it into loaves.

This project represents our preliminary efforts to develop an objective and realistic methodology for assessing the capabilities of AI systems to autonomously solve research-level math questions. After letting these ideas ferment in the community, we hope to produce a more structured benchmark.

We present a diverse set of 10 research-level math questions, drawn from algebraic combinatorics, spectral graph theory, algebraic topology, stochastic analysis, symplectic geometry, representation theory, lattices in Lie groups, tensor analysis, and numerical linear algebra. Each question arose naturally in the research process of the authors and has been answered with a proof of roughly five pages or less, but the answers have not yet been posted online.

For the next batch, we will implement a benchmarking phase prior to the community release. The benchmark phase will be designed to ensure the following features:

- Verification that the solutions are produced autonomously.
- A formal grading scheme and refereeing, modeled on the journal review system.
- An explicit description of the problem selection process, including advance internal testing on systems which have a zero data retention policy.

If you are interested in an assessment of your solutions to the next round of questions, please email contact@1stproof.org.

<https://1stproof.org/>
<https://codeberg.org/tgkolda/1stproof/src/branch/main/2026-02-batch/FirstProofSolutionsComments.pdf>

First Proof

Mohammed Abouzaid*
Stanford University

Andrew J. Blumberg
Columbia University

Martin Hairer
EPFL and Imperial

Joe Kileel
University of Texas at Austin

Tamara G. Kolda
MathSci.ai

Paul D. Nelson
Aarhus University

Daniel Spielman
Yale University

Nikhil Srivastava†
University of California, Berkeley

Rachel Ward‡
University of Texas at Austin

Shmuel Weinberger
University of Chicago

Lauren Williams§
Harvard University

February 6, 2026

Feb 24, 2026

44

22

Solving Unseen Mathematical Problems

3 The questions

- Let $\mathbb{T}^3$ be the three dimensional unit size torus and let μ be the Φ_3^2 measure on the space of distributions $\mathcal{D}'(\mathbb{T}^3)$. Let $\psi: \mathbb{T}^3 \rightarrow \mathbb{R}$ be a smooth function that is not identically zero and let $T_\psi: \mathcal{D}'(\mathbb{T}^3) \rightarrow \mathcal{D}'(\mathbb{T}^3)$ be the shift map given by $T_\psi(u) = u + \psi$ (with the usual identification of smooth functions as distributions). Are the measures μ and $T_\psi\mu$ equivalent? Here, equivalence of measures is in the sense of having the same null sets and T_ψ denotes the pushforward under T_ψ .
- Let F be a non-archimedean local field with ring of integers $\mathfrak{o}$. Let N_π denote the subgroup of $\mathrm{GL}_n(F)$ consisting of upper-triangular unipotent elements. Let $\psi: F \rightarrow \mathbb{C}^\times$ be a nontrivial additive character of conductor $\mathfrak{o}$, identified in the standard way with a generic character of N_π . Let Π be a generic irreducible admissible representation of $\mathrm{GL}_{n+1}(F)$, realized in its ψ^{-1} -Whittaker model $\mathcal{W}(\Pi, \psi^{-1})$. Must there exist $W \in \mathcal{W}(\Pi, \psi^{-1})$ with the following property?
Let π be a generic irreducible admissible representation of $\mathrm{GL}_n(F)$, realized in its ψ -Whittaker model $\mathcal{W}(\pi, \psi)$. Let $\mathfrak{q}$ denote the conductor ideal of π , let $Q \in F^\times$ be a generator of $\mathfrak{q}^{-1}$, and set

$$w_Q := I_{n+1} + Q E_{n,n+1} \in \mathrm{GL}_{n+1}(F),$$

where E_{ij} is the matrix with a 1 in the (i, j) -entry and 0 elsewhere. For some $V \in \mathcal{W}(\pi, \psi)$, the local Rankin-Selberg integral

$$\int_{N_{\pi} \backslash \mathrm{GL}_{n+1}(F)} W(\mathrm{diag}(g, 1)w_Q) V(g) |\det g|^{-\frac{1}{2}} dg$$

is finite and nonzero for all $s \in \mathbb{C}$.

4 Comments on solutions

On February 4 and 5, 2026, we tested the questions on Gemini 3.0 Deep Think and ChatGPT 5.2 Pro with the following prompts.

Prompt 1. The following is a research-level math question. The question has an answer, but it might not appear on the internet. Please make a best effort to provide a rigorous and complete answer to the question. Write the output as a compilable LaTeX document using the standards of rigor and scholarship that prevail in the mathematical literature.

Prompt 2 (Internet Discouraged). The following is a research-level math question. The question has an answer. Please make a best effort to provide a rigorous and complete answer to the question. Write the output as a compilable LaTeX document using the standards of rigor and scholarship that prevail in the mathematical literature. Do not use web search, but instead try to reason through the answer.

<https://1stproof.org/>
<https://codeberg.org/tqkolda/1stproof/src/branch/main/2026-02-batch/FirstProofSolutionsComments.pdf>

4.1 Question 1: Martin Hairer

In this case, a note with a very short sketch of proof (far short of the level of detail one would expect for a published article) was posted on the author's homepage some time ago. The answer given by GPT-Pro simply quotes that note, claiming that it contains a detailed proof of the result. This is incorrect and it is despite the LLM being specifically instructed to comply with "mathematics publication" levels of scholarship. (Taking for granted a result that is merely stated in an unpublished note with a very rough sketch of proof is not considered acceptable in the mathematics literature.)

Another behaviour we observed was that the LLM would take as a premise the (wrong!) statement that the Φ_3^2 measure is equivalent to the free field measure, from which it then

6

correctly deduces the (incorrect) claim that the Φ_3^2 measure is quasi-invariant under smooth shifts.

4.6 Question 6: Daniel Spielman

Gemini asserted that it presented a proof of the existence of a constant that satisfied Question 6. But, after some correct statements, it presented a very vague explanation of how the proof could be finished. To me, it seems unlikely that the approach can be turned into a correct proof.


ChatGPT 5.2 Pro asserted that it could not answer the question. So, it instead offered a correct upper bound of 1/2 on the constant, if it exists.

Feb 24, 2026

45

The Anthropic Hive Mind

Steve Yegge Follow 16 min read · 4 days ago




Anthropic Hive Mind:

- **Radical Transparency**
 - Death of ego
- **Extreme Velocity**
 - Claude Cowork launched in 10 days.
- **"Yes, and..." Improvisational Style**
 - Every idea is welcomed, examined, and judged by the collective.

<https://x.com/MrinankSharma/status/2020881722003583421>; <https://steve-yegge.medium.com/the-anthropic-hive-mind-d01f768f3d7b>

Feb 24, 2026

2026, Best of Time, Worst of Time



46